

微信公众号与哔哩哔哩有视频版本同步：

1.

<https://mp.weixin.qq.com/s?biz=Mzk0OTI4NzgyNg==&tempkey=MTEzM19GVHpxWTdsbFRTWDJPQjNSZjgySXUxSDRwekVUbnhHdzFOR0tQWGPvYj1PSW5ULTJZcHRGT2pMaERBUVVUX0s2VX1STVQ0c2VXS WN5TjZCa0ZNNtctYm91Xy1kY1dBTENkeWtLZ2p3UE5UcTFuUzJrYj1MZH0czVJYUd0LV1acjdBRzVOYUtvMFVQZXRZmxneHJMUjBJSi1zY1d2eEdnU0t6ZDhBfn4%3D&chksm=435beff3742c66e58e0ac2342fa14426eb498ad1aa6c46e4d86612b60de0d9684cd50597db18#rd>

Page 1

## Motivation

- Various model compression methods rely on the assumption that deep networks are over-parametrized
- This redundancy in deep networks can appear *explicitly* as well as *implicitly*
- Several structured and unstructured pruning methods exploit the explicit redundancy by removing redundant components in the network (filter-level or parameter-level)
- Tensor decomposition methods (e.g. WeightSVD, SpatialSVD) remove somewhat implicit elements, i.e. singular values, to construct low-rank decompositions of tensors for efficient inference
- Redundancy can also be seen as model weights possessing unnecessarily high *degrees of freedom*. Some works efficient basis representations to reduce the DOF by decreasing number of trainable parameters [14, 29, 38]
  - Systematically choosing a basis can lead to efficient benefits during inference as well, not just training

### Contributions

- In this work, we propose a neat way to restrict the DOFs of weight tensors, which eventually leads to a decomposition of the convolution operation into efficient, scaled, sum-pooling components
- We also propose Structural Regularization, an effective training method to enable this so-called structural decomposition with minimal loss in accuracy


Page2

# Composite Kernels

**Definition 1:** For  $\mathbb{R}^{C \times N \times N}$ , a Composite Basis  $\mathbb{B} = \{\beta_1, \beta_2, \dots, \beta_M\}$  is a linearly independent set of binary tensors of dimension  $C \times N \times N$  as its basis elements. That is,  $\beta_m \in \mathbb{R}^{C \times N \times N}$ , each element  $\beta_{mijk}$  of  $\beta_m \in \{0,1\}$  and  $\sum_{m=1}^M \alpha_m \beta_m = 0$  iff  $\alpha_m = 0 \forall m$ .

**Definition 2:** A kernel  $W \in \mathbb{R}^{C \times N \times N}$  is a Composite Kernel if it lies in the subspace of a Composite Basis  $\mathbb{B}$ , i.e. it can be constructed as a linear combination of the elements of  $\mathbb{B}$ :  $\exists \alpha = [\alpha_1, \dots, \alpha_M]$  such that  $W := \sum_{m=1}^M \alpha_m \beta_m$

Example below shows A  $3 \times 3$  Composite Kernel with different underlying structures / basis. Left two kernels possess 4 DOF, whereas the rightmost kernel possesses 6 DOF.



Page3

## Convolution with a Composite Kernels

- Consider a convolution operation with a Composite Kernel  $W$  of size  $C \times N \times N$ . To compute an output,  $W$  is convolved with a  $C \times N \times N$  volume of the input feature map

$$X * W = X * \sum_{m=1}^M \alpha_m \beta_m = \sum_{m=1}^M \alpha_m (X * \beta_m) = \sum_{m=1}^M \alpha_m \text{sum}(X \bullet \beta_m) = \sum_{m=1}^M \alpha_m E_m \quad (1)$$

- Ordinarily, the convolution  $X * W$  would involve  $CN^2$  multiplications and  $CN^2 - 1$  additions
- But since  $\beta_m$  is a binary tensor, the operation  $\text{sum}(X \bullet \beta_m)$  is same as adding the elements of  $X$  wherever  $\beta_m = 1$ , thus no multiplications!
- Hence, from (1), we see that we only need  $M$  multiplications, and the number of additions are given by:

$$\text{Num additions} = \sum_{m=1}^M \underbrace{(\text{sum}(\beta_m) - 1)}_{\text{from } \text{sum}(X \bullet \beta_m)} + \underbrace{(M - 1)}_{\text{from } \sum \alpha_m E_m} = \sum_{m=1}^M \text{sum}(\beta_m) - 1 \quad (2)$$


- Depending on the structure ( $\mathbb{B}$ ), num additions *per output* can be larger than  $CN^2 - 1$ , e.g. in the second example on previous slide where  $C = 1, N = 3, M = 4$ , we have  $\sum_m \text{sum}(\beta_m) - 1 = 15 > CN^2 - 1$
- Next, we show that the num additions can be amortized to as low as  $M - 1$  ☺

Pag5


# Structured Kernels

**Definition 3:** A kernel in  $\mathbb{R}^{C \times N \times N}$  is a Structured Kernel if it is a Composite Kernel  $M = cn^2$  for some  $1 \leq n \leq N$  and  $1 \leq c \leq C$ , and if each basis tensor  $\beta_m$  is made of a  $(C - c + 1) \times (N - n + 1) \times (N - n + 1)$  cuboid of 1's, while rest of its coefficients being 0.

- A Structured Kernel is characterized by its dimensions  $C, N$  and its underlying parameters  $c, n$ .
- Convolutions performed using Structured Kernels are called Structured Convolutions.



This is a 2D case of a  $3 \times 3$  structured kernel where  $C = 1$ ,  $N = 3$ ,  $c = 1$ ,  $n = 2$ . As shown, there are  $M = cn^2 = 4$  basis elements and each element has a  $2 \times 2$  sized patch of 1's



This shows a 3D case where  $C = 4$ ,  $N = 3$ ,  $c = 2$ ,  $n = 2$ . Here, there are  $M = cn^2 = 8$  basis elements and each element has a  $2 \times 2 \times 2$  cuboid of 1's. Note how these cuboids of 1's (shown in colors) cover the entire  $4 \times 3 \times 3$  grid

Page6

## Structured Convolutions

- A major advantage of Structured Kernels is that all the basis elements are just shifted versions of each other
- This means, considering the convolution  $X * W$  (in Eq. (1)) for the *entire* feature map  $X$ , the summed outputs  $X * \beta_m$  for all  $\beta_m$ s **are the same** (except on the edges of  $X$ )
- As a result, the outputs  $\{X * \beta_1, \dots, X * \beta_{cn^2}\}$  can be computed using a single sum-pooling operation on  $X$  with a kernel size of  $(C - c + 1) \times (N - n + 1) \times (N - n + 1)$
- Example below shows how a convolution with a  $3 \times 3$  structured kernel can be broken into a  $2 \times 2$  sum-pooling followed by a  $2 \times 2$  convolution with a kernel made of  $\alpha_i$ 's:


$$\begin{array}{c}
 \begin{array}{|c|c|c|} \hline x_1 & x_2 & x_3 \\ \hline x_4 & x_5 & x_6 \\ \hline x_7 & x_8 & x_9 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline \alpha_1 & \alpha_1 + \alpha_2 & \alpha_2 \\ \hline \alpha_1 + \alpha_3 & \alpha_1 + \alpha_2 + \alpha_3 + \alpha_4 & \alpha_2 + \alpha_4 \\ \hline \alpha_3 & \alpha_3 + \alpha_4 & \alpha_4 \\ \hline \end{array} = \begin{array}{l} \alpha_1 * (x_1 + x_2 + x_4 + x_5) \\ + \alpha_2 * (x_2 + x_3 + x_5 + x_6) \\ + \alpha_3 * (x_4 + x_5 + x_7 + x_8) \\ + \alpha_4 * (x_5 + x_6 + x_8 + x_9) \end{array} \\
 \text{Convolution with } 3 \times 3 \text{ kernel} \\
 \hline
 = \begin{array}{|c|c|c|} \hline x_1 & x_2 & x_3 \\ \hline x_4 & x_5 & x_6 \\ \hline x_7 & x_8 & x_9 \\ \hline \end{array} \xrightarrow{\text{Sum Pooling}} \begin{array}{|c|c|} \hline x_1 + x_2 + x_4 + x_5 & x_2 + x_3 + x_5 + x_6 \\ \hline x_4 + x_5 + x_7 + x_8 & x_5 + x_6 + x_8 + x_9 \\ \hline \end{array} * \begin{array}{|c|c|} \hline \alpha_1 & \alpha_2 \\ \hline \alpha_3 & \alpha_4 \\ \hline \end{array} \\
 \text{2} \times \text{2 sum-pooling + convolution with 2} \times \text{2 kernel}
 \end{array}$$

Page4



# Structural Decomposition

- Consider a general convolution layer of size  $C_{out} \times C \times N \times N$  that has  $C_{out}$  kernels of size  $C \times N \times N$
- In the proposed design, the same underlying basis  $\mathbb{B} = \{\beta_1, \dots, \beta_{cn^2}\}$  is used for all  $C_{out}$  kernels in the layer
- Suppose any two structured kernels in this layer,  $W_1 := \sum_{m=1}^M \alpha_m^{(1)} \beta_m$  and  $W_2 := \sum_{m=1}^M \alpha_m^{(2)} \beta_m$ . The convolution outputs will be,  $X * W_1 = \sum_m \alpha_m^{(1)} (X * \beta_m)$  and  $X * W_2 = \sum_m \alpha_m^{(2)} (X * \beta_m)$
- The computation  $X * \beta_m$  is common to all the kernels of this layer. **Hence, the sum-pooling only needs to be computed once and reused across all layers.** This important result leads to the following decomposition:



**Figure:** Decomposition of Structured Convolution. On the left, the conventional operation of a structured convolutional layer of size  $C_{out} \times C \times N \times N$  is shown. On the right, we show that it is equivalent to performing a 3D sum-pooling followed by a convolutional layer of size  $C_{out} \times c \times n \times n$ .

## Improvement in Computational Complexity

- The following table shows a comparison of the number of parameters and operations (multiplications and additions) before and after the structural decomposition

	Conventional convolution (before decomposition)	Sum-pooling + Smaller convolution (after decomposition)
# Parameters	$C_{out} C N^2$	$C_{out} c n^2$
# Multiplications	$C N^2 \times C_{out} H' W'$	$c n^2 \times C_{out} H' W'$
# Additions	$(C N^2 - 1) \times C_{out} H' W'$	$((C - D + 1)(N - k + 1)^2 - 1) \times c H_1 W_1 + (c n^2 - 1) \times C_{out} H' W'$

- We see that both number of parameters and multiplications have reduced by a factor of  $cn^2/CN^2$
- Number of additions after decomposition can be re-written as:

$$C_{out} \left( \frac{((C - c + 1)(N - n + 1)^2 - 1) c H_1 W_1}{C_{out}} + (c n^2 - 1) H' W' \right)$$

- Hence, when  $C_{out}$  is large enough, the first term inside parentheses (the sum-pooling component) gets amortized and num additions become  $\approx (c n^2 - 1) \times C_{out} H' W' = (M - 1) \times C_{out} H' W' \dots$  since  $M = c n^2$ 
  - Hence, additions also decrease by approximately  $cn^2/CN^2$

## Applicability of Structural Decomposition

- Standard convolution ( $C \times N \times N$ ), depthwise convolution ( $1 \times N \times N$ ), and pointwise convolution ( $C \times 1 \times 1$ ) kernels can all be constructed as 3D structured kernels, which means that this decomposition can be widely applied to existing architectures
- It can also be extended to Fully Connected layers by noting that the linear operation  $WX$  is mathematically equivalent to a  $1 \times 1$  convolution  $\text{unsqueezed}(X) * \text{unsqueezed}(W)$ , where  $\text{unsqueezed}(X)$  is the same as  $X$  but with dimensions  $Q \times 1 \times 1$  and  $\text{unsqueezed}(W)$  is the same as  $W$  but with dimensions  $P \times Q \times 1 \times 1$



Figure: Structural decomposition of a matrix multiplication

- Same as before, we get a reduction in both number of parameters and multiplications by a factor of  $R/Q$  and number of additions by a factor of  $\frac{R(Q-R)+(PR-1)P}{(PQ-1)P}$

## Structural Regularization

- We propose training a deep network with a Structural Regularization loss that can gradually push the deep network's kernels to be structured via training:

$$\mathcal{L}_{total} = \mathcal{L}_{task} + \lambda \sum_{l=1}^L \frac{\|(I - A_l A_l^+) W_l\|_F}{\|W_l\|_F}$$

### Proposed Training Scheme:

- Step 1: Train the original architecture with the Structural Regularization loss.
  - After Step 1, all weight tensors in the deep network will be *almost* structured.
- Step 2: Apply the decomposition on every layer and compute  $\alpha_l = A_l^+ W_l$ .
  - This results in a smaller and more efficient decomposed architecture with  $\alpha_l$ 's as the weights. Note that, every convolution / linear layer from the original architecture is now replaced with a sum-pooling layer and a smaller convolution / linear layer.
- We show in the paper that the Structural Regularization training method leads to significantly better performance than just training the Structured Convolutions as an architectural feature

# Results – Image Classification with ResNets

**Table: ResNets on CIFAR-10**

Architecture	Adds ( $\times 10^9$ )	Mults ( $\times 10^6$ )	Params ( $\times 10^6$ )	Acc. (in %)
ResNet56	125.49	126.02	0.85	93.03
<b>Struct-56-A (ours)</b>	63.04	62.10	0.40	92.65
<b>Struct-56-B (ours)</b>	48.49	33.87	0.21	91.78
Ghost-Res56 [7]	63.00	63.00	0.43	92.70
ShiftRes56-6 [42]	51.00	51.00	0.58	92.69
AMC-Res56 [11]	63.00	63.00	–	91.90
ResNet32	68.86	69.16	0.46	92.49
<b>Struct-32-A (ours)</b>	35.59	35.09	0.22	92.07
<b>Struct-32-B (ours)</b>	26.29	18.18	0.11	90.24
ResNet20	40.55	40.74	0.27	91.25
<b>Struct-20-A (ours)</b>	20.77	20.42	0.13	91.04
<b>Struct-20-B (ours)</b>	15.31	10.59	0.067	88.47
ShiftRes20-6 [42]	23.00	23.00	0.19	90.59

**Table: ResNets on ImageNet**

Architecture	Adds ( $\times 10^9$ )	Mults ( $\times 10^9$ )	Params ( $\times 10^6$ )	Acc. (in %)
ResNet50	4.09	4.10	25.56	76.15
<b>Struct-50-A (ours)</b>	2.69	2.19	13.49	75.65
<b>Struct-50-B (ours)</b>	1.92	1.38	8.57	73.41
ChPrune-R50-2x [12]	2.04	2.05	17.89	75.44
WeightSVD-R50 [48]	2.04	2.05	13.40	75.12
Ghost-R50 (s=2) [7]	2.20	2.20	13.00	75.00
Versatile-v2-R50 [40]	3.00	3.00	11.00	74.50
ShiftResNet50 [42]	–	–	11.00	73.70
Slim-R50 0.5x [45]	1.10	1.10	6.90	72.10
ResNet34	3.66	3.67	21.80	73.30
<b>Struct-34-A (ours)</b>	1.71	1.71	9.82	72.81
<b>Struct-34-B (ours)</b>	1.47	1.11	5.60	69.44
ResNet18	1.81	1.82	11.69	69.76
<b>Struct-18-A (ours)</b>	0.88	0.89	5.59	69.13
<b>Struct-18-B (ours)</b>	0.79	0.63	3.19	66.19
WeightSVD-R18 [48]	0.89	0.90	5.83	67.71
ChPrune-R18-2x [12]	0.90	0.90	7.45	67.69
ChPrune-R18-4x	0.44	0.45	3.69	61.56

**Table: MobileNetV2 and EfficientNet**

Architecture	Adds ( $\times 10^6$ )	Mults ( $\times 10^6$ )	Params ( $\times 10^6$ )	Acc. (in %)
MobileNetV2	0.30	0.31	3.50	72.19
<b>Struct-V2-A (ours)</b>	0.26	0.23	2.62	71.29
<b>Struct-V2-B (ours)</b>	0.29	0.17	1.77	64.93
AMC-MV2 [11]	0.21	0.21	–	70.80
ChPrune-MV2-1.3x	0.22	0.23	2.58	68.99
Slim-MV2 [45]	0.21	0.21	2.60	68.90
WeightSVD 1.3x	0.22	0.23	2.45	67.48
ChPrune-MV2-2x	0.15	0.15	1.99	63.82
EfficientNet-B1 [37]	0.70	0.70	7.80	78.50 <sup>1</sup>
<b>Struct-EffNet (ours)</b>	0.62	0.45	4.93	76.40 <sup>1</sup>
EfficientNet-B0 [37]	0.39	0.39	5.30	76.10 <sup>1</sup>

## Concluding remarks

- In this work, we proposed Composite Kernels and Structured Convolutions in an attempt to exploit redundancy in the implicit structure of convolution kernels
- We show that Structured Convolutions can be decomposed into a computationally cheap sum-pooling component followed by a significantly smaller convolution, by training the model using a Structural Regularization loss.
- Sum-pooling relies purely on additions, which are known to be extremely power-efficient. **Hence, our method shows promise in deploying deep models on low-power devices.**
- Since our method keeps the convolutional structures, it allows integration of further model compression schemes, which we leave as future work.